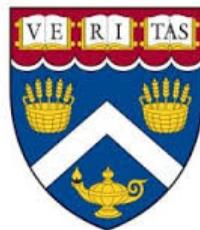


Final Project

# Object Classification of Photo Instance Segmentation

Abderrazzaq, Saad



CSCI E-89 Deep Learning  
**Harvard University Extension School**  
Prof. Zoran B. Djordjević

# Problem Statement

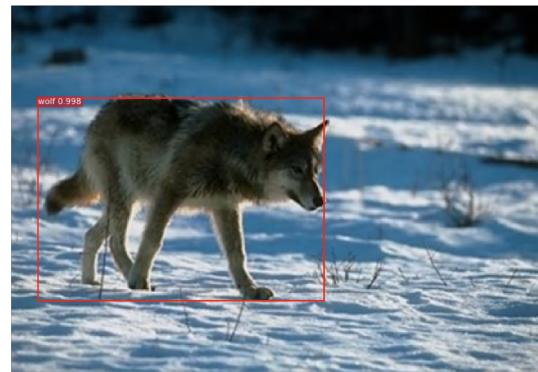
- To identify instances in an image using instance segmentation by using an open source library that implements a method called Mask R-CNN on top of Keras/Tensorflow.
- Demonstrate single and multi-class instance segmentation as well as the process of annotating data. I use data from ImageNet (wolfs, signs) and my own photos to train and test image segmentation. The data is annotated using VGG Image Annotator (VIA).

# Instance Segmentation

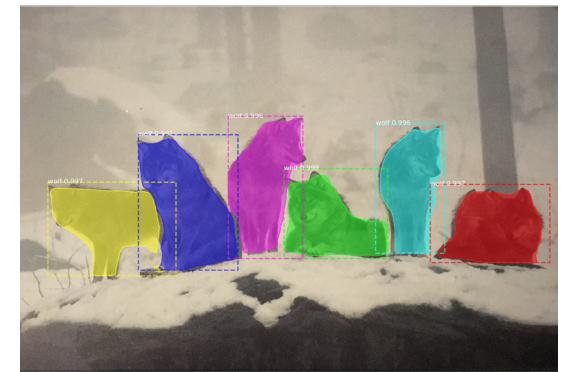
Image  
Classification



Object  
Detection



Instance  
Segmentation



- Detection of one or more objects of different classes
- Differentiate instances
- Label each pixel with a category label
- Generate segment mask for each instance

# Technology

- Anaconda
  - Python (Version 3.6.4)
  - Jupyter Notebook
  - Packages: NumPy, SciPy, Matplotlib, etc
- Keras (v2.1.5) / Tensorflow (v1.6.0)
- OpenCV 3.4 (i.e. cv2)
- Mask\_RCNN ([https://github.com/matterport/Mask\\_RCNN.git](https://github.com/matterport/Mask_RCNN.git))
- pycocotools
- Imgaug (<https://github.com/aleju/imgaug>)
- VGG Image Annotator (v1.0.5)

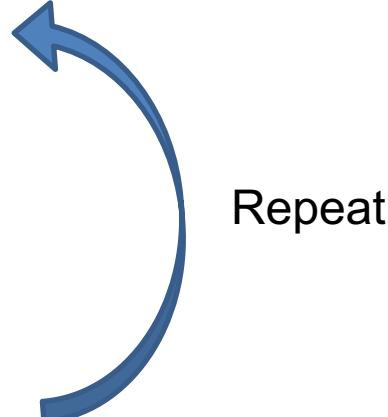
# Data Set

- Dataset consists of images download from ImageNet
- The images came in a variety of sizes and quality and were all RGB.
- The first data batch consistent of variety of wolf species.
  - Out of 1391 images only 171 images annotated.
  - The total size of this data was 15M.
- The second batch consistent of traffic signs
  - Used to train the network to detect multiple traffic signs.
  - This batch consisted of 85 annotated images
  - Total size was 26M
- The VGG Image Annotator (VIA) was used to annotate the images.

# Installation Overview

1. Download Anaconda
2. Download and install Ubuntu pre-requisites for OpenCV
3. Download and build OpenCV
4. Install opencv-python
5. Install Cython
6. Git clone pycocotools
7. Git clone imgaug
8. Git clone mask\_rcnn
9. Download and unzip VGG Image Annotator

# Process

1. Create Configuration and DataSet code for:
    1. Loading images. (mrcnn has good examples)
    2. Loading image masks
    3. Splitting images and annotations between training and validation
  2. Annotate Images
  3. Validate data
  4. Run training
  5. Run detection tests
  6. Debug/analyze
- 
- Repeat

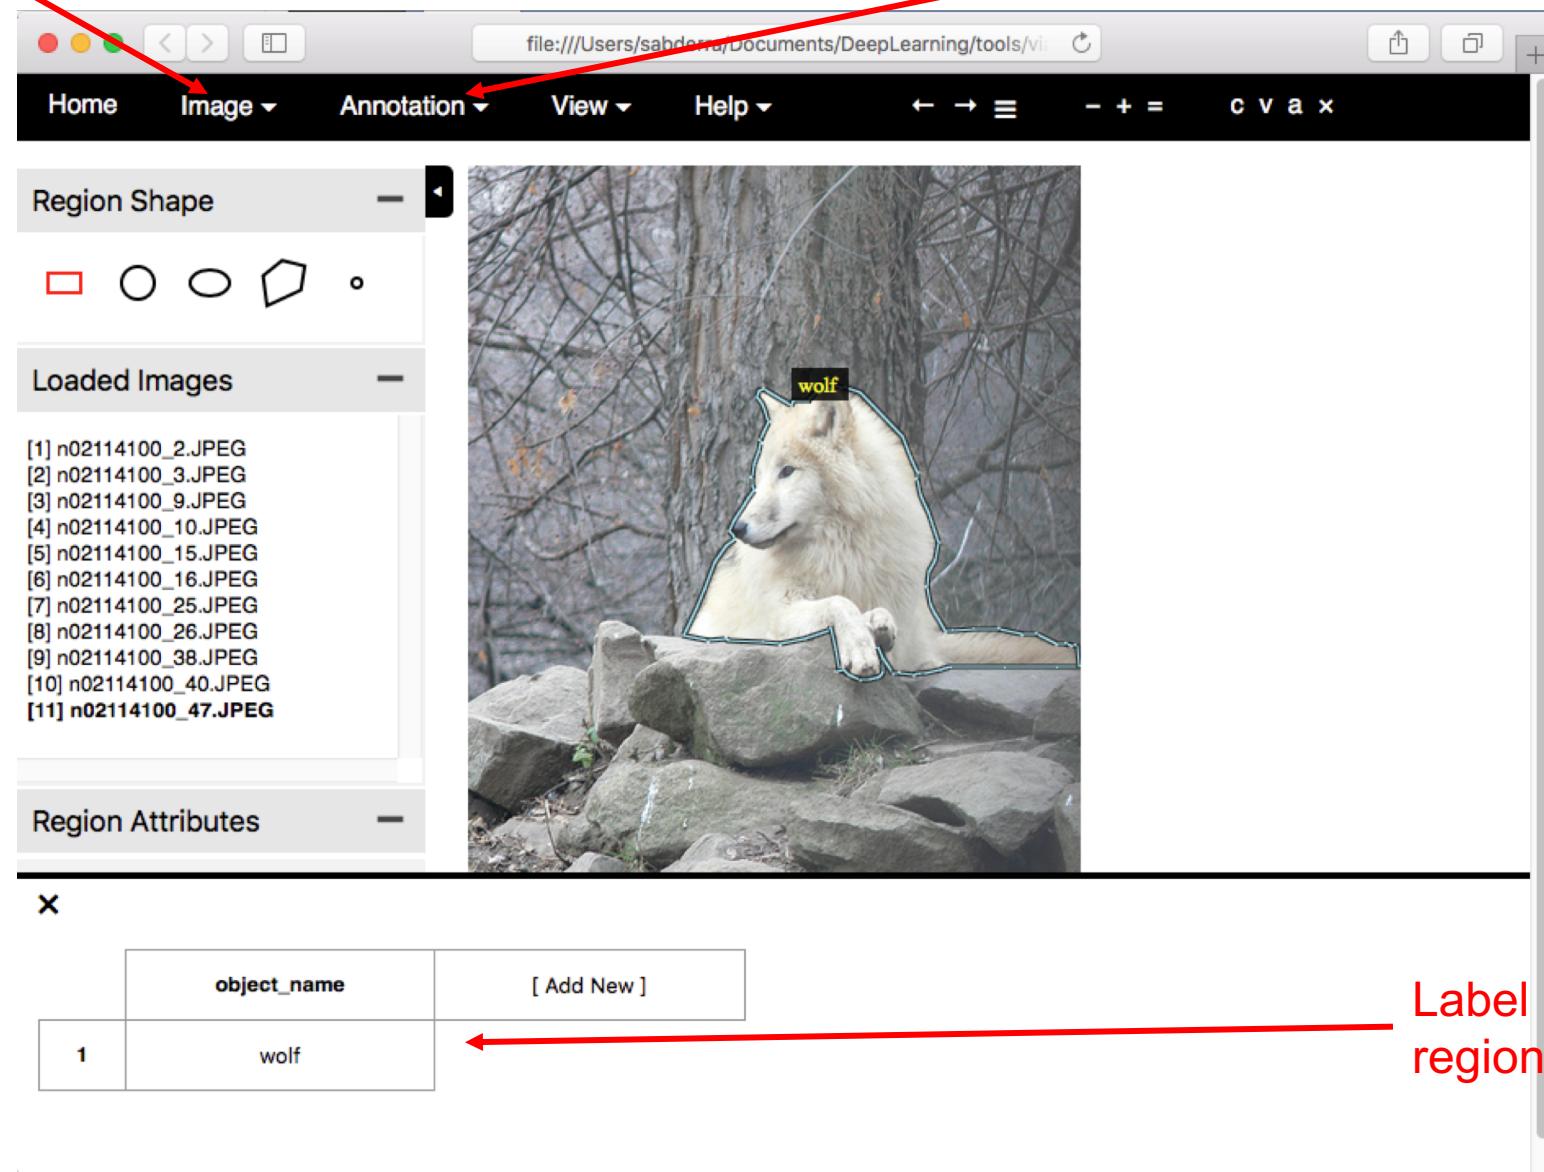
# Concepts

- Region Proposals
  - An algorithm that identifies those rectangular regions giving them an ‘objectness’ score. This can be up to 1000 regions. The region of interest (ROI) is further operated on and filtered to remove all by the highest confidence of containing an object matching one of the classes.
- Intersection Over Union (IoU)
  - A measure used when evaluating object localization. This helps to narrow down the bounding boxes to the one best fitting the object. Uses the ground truth bounding box and the predicted bounding box. IoU is area of overlap over area of union.
- Non Max Suppression
- Feature Pyramids

Load Images

# VGG Image Annotator (VIA)

Import/Save Annotations



Label each region

# VIA JSON Format

- VIA exports annotations in both csv and json formats. I used json

```
{  
  "n02114100_10.JPEG81442": {  
    "fileref": "",  
    "size": 81442,  
    "filename": "n02114100_10.JPG",  
    "base64_img_data": "",  
    "file_attributes": {},  
    "regions": {  
      "0": {  
        "shape_attributes": {  
          "name": "polygon",  
          "all_points_x": [ ... ],  
          "all_points_y": [ ... ]  
        },  
        "region_attributes": {  
          "object_name": "wolf"  
        }  
      }  
    }  
  },  
  "n02114100_120.JPEG174337": {  
    ...  
  },  
  "n02114100_123.JPG4410": {  
    ...  
  },  
  ...  
}
```

- The key attributes are:
  - filename is the location of the image on disk
  - all\_points\_x and all\_points\_y define the polygons for the masks and deriving the bounding boxes.
  - object\_name was used to label the specific polygon. This needs to be mapped to the class index.
  - Note, shape of the image is missing and not written by via v1.0.5.

# Dataset - load\_by\_annotations

```
17 def load_by_annotations(self, dataset_dir, annotations_list, class_ids):
18     """Load a specific set of annotations and from them images.
19     dataset_dir: Root directory of the dataset.
20     annotations_list: The annotations (and images) to be loaded.
21     class_ids: List of classes to use.
22     """
23
24     ann_class_names = set()
25     for a in annotations_list:
26         regions = a['regions']
27         for r, v in regions.items():
28             ann_class_names.add(v['region_attributes']['object_name'])
29
30     # Add classes.
31     for i, name in enumerate(ann_class_names):
32         # Skip over background if it appears in the class name list
33         index = i + 1
34         if name != 'BG':
35             print('Adding class {:3}:{}'.format(index, name))
36             self.add_class('sign', index, name)
37             self.map_name_to_id[name] = index
38
39     # Add images
40     for a in annotations_list:
41         # Get the x, y coordinates of points of the polygons that make up
42         # the outline of each object instance. There are stores in the
43         # shape_attributes (see json format above)
44         polygons = [r['shape_attributes'] for r in a['regions'].values()]
45
46         r_object_name = [r['region_attributes']['object_name'] for r in a['regions']]
47
48         assert len(polygons) == len(r_object_name)
49
50         # load_mask() needs the image shape. This is only manageable since
51         # the dataset is small. Consider pre-processing and/or updating via.
52         image_path = os.path.join(dataset_dir, a['filename'])
53         image = skimage.io.imread(image_path)
54         height, width = image.shape[:2]
55
56         self.add_image(
57             "sign",
58             image_id=a['filename'], # use file name as a unique image id
59             path=image_path,
60             width=width, height=height,
61             polygons=polygons,
62             r_object_name=r_object_name)
63
```

- Given a list of annotations, this function loads the classes, polygons, and labels.
- Annotations are collected by the split\_annotations() fcn

# Dataset - load\_mask

```
67 def load_mask(self, image_id):
68     """Generate instance masks for an image.
69     Returns:
70     masks: A bool array of shape [height, width, instance count] with
71         one mask per instance.
72     class_ids: a 1D array of class IDs of the instance masks.
73     """
74
75     # If not an object in our dataset image, delegate to parent class.
76     image_info = self.image_info[image_id]
77     if image_info["source"] not in self.class_names:
78         print("warning: source {} not part of our classes, delegating to parent")
79         return super(self.__class__, self).load_mask(image_id)
80
81     # Convert polygons to a bitmap mask of shape
82     # [height, width, instance_count]
83     info = self.image_info[image_id]
84     mask = np.zeros([info["height"], info["width"], len(info["polygons"])],
85                    dtype=np.uint8)
86     class_ids = []
87     for i, p in enumerate(info["polygons"]):
88         # Get indexes of pixels inside the polygon and set them to 1
89         try:
90             rr, cc = skimage.draw.polygon(p['all_points_y'], p['all_points_x'])
91             mask[rr, cc, i] = 1
92             class_id = self.map_name_to_id[info['r_object_name'][i]]
93             class_ids.append(class_id)
94         except:
95             print(image_info)
96             raise
97
98     # Return mask, and array of class IDs of each instance. Since we have
99     # one class ID only, we return an array of 1s
100    class_ids = np.array(class_ids, dtype=np.int32)
101
102    return mask.astype(np.bool), class_ids
```

- Overridden from Mask RCNN *load\_mask* to define how to load from the via format.
- *map\_name\_to\_id* is used for associating the label string with the image id int.

# Image Augmentation

```
11 import imgaug as ia
12 from imgaug import augmenters as iaa
13
14 ia.seed(1)
15
16 # http://imgaug.readthedocs.io/en/latest/source/augmenters.html#sequential
17 seq_of_aug = iaa.Sequential([
18     iaa.Crop(percent=(0, 0.1)), # random crops
19
20     # horizontally flip 50% of the images
21     # iaa.Fliplr(0.5), # Does not make sense for signs
22
23     # Gaussian blur to 50% of the images
24     # with random sigma between 0 and 0.5.
25     iaa.Sometimes(0.5,
26         iaa.GaussianBlur(sigma=(0, 0.5))
27     ),
28
29     # Strengthen or weaken the contrast in each image.
30     iaa.ContrastNormalization((0.75, 1.5)),
31
32     # Add gaussian noise.
33     # For 50% of all images, we sample the noise once per pixel.
34     # For the other 50% of all images, we sample the noise per pixel AND
35     # channel. This can change the color (not only brightness) of the
36     # pixels.
37     iaa.AdditiveGaussianNoise(loc=0, scale=(0.0, 0.05*255), per_channel=0.5),
38
39     # Make some images brighter and some darker.
40     # In 20% of all cases, we sample the multiplier once per channel,
41     # which can end up changing the color of the images.
42     iaa.Multiply((0.8, 1.2), per_channel=0.2),
43
44     # Apply affine transformations to each image.
45     # Scale/zoom them from 90% to 110%
46     # Translate/move them, rotate them
47     # Shear them slightly -2 to 2 degrees.
48     iaa.Affine(
49         scale={"x": (0.9, 1.1), "y": (0.9, 1.1)},
50         translate_percent={"x": (-0.2, 0.2), "y": (-0.2, 0.2)},
51         rotate=(-5, 5),
52         shear=(-2, 2)
53     )
54 ], random_order=True) # apply augmenters in random order
```

- The imgaug library is used
- Supports deterministic setting for image/mask
- Need to test augmenter results to ensure safe

# Training

```
14 model.load_weights(COCO_MODEL_PATH, by_name=True, exclude=[  
15     "mrcnn_class_logits", "mrcnn_bbox_fc",  
16     "mrcnn_bbox", "mrcnn_mask"])
```

```
1 def train(model, dataset_path, epochs=30):  
2     """Train the model."""  
3  
4     # Create the train and val dataset.  
5     dataset_train, dataset_val = create_datasets(dataset_path+'/train', config.CLASS_NAMES)  
6  
7     # Prepare them  
8     dataset_train.prepare()  
9     dataset_val.prepare()  
10  
11    # Experiment with training options.  
12    # Since we're using a very small dataset, and starting from  
13    # COCO trained weights, we don't need to train too long. Also,  
14    # no need to train all layers, just the heads should do it.  
15    print("Training network heads")  
16    history = model.train(dataset_train, dataset_val,  
17                          learning_rate=config.LEARNING_RATE,  
18                          epochs=epochs,  
19                          layers='heads',  
20                          augmentation=seq_of_aug  
21                          )  
22  
23    return history
```

```
1 history = train(model, dataset_path, 75)
```

Use a pre trained model for training

Split the data between training and val  
Default is .8 for training.

Wrapping from MaskRCNN that builds and compiles a model using Keras

Kick off the training

# Training Parameters

- No significant tuning was performed. I attempted to tune learning rate, activation, batch normalization, but changes did not out do default config.
- Configuration:

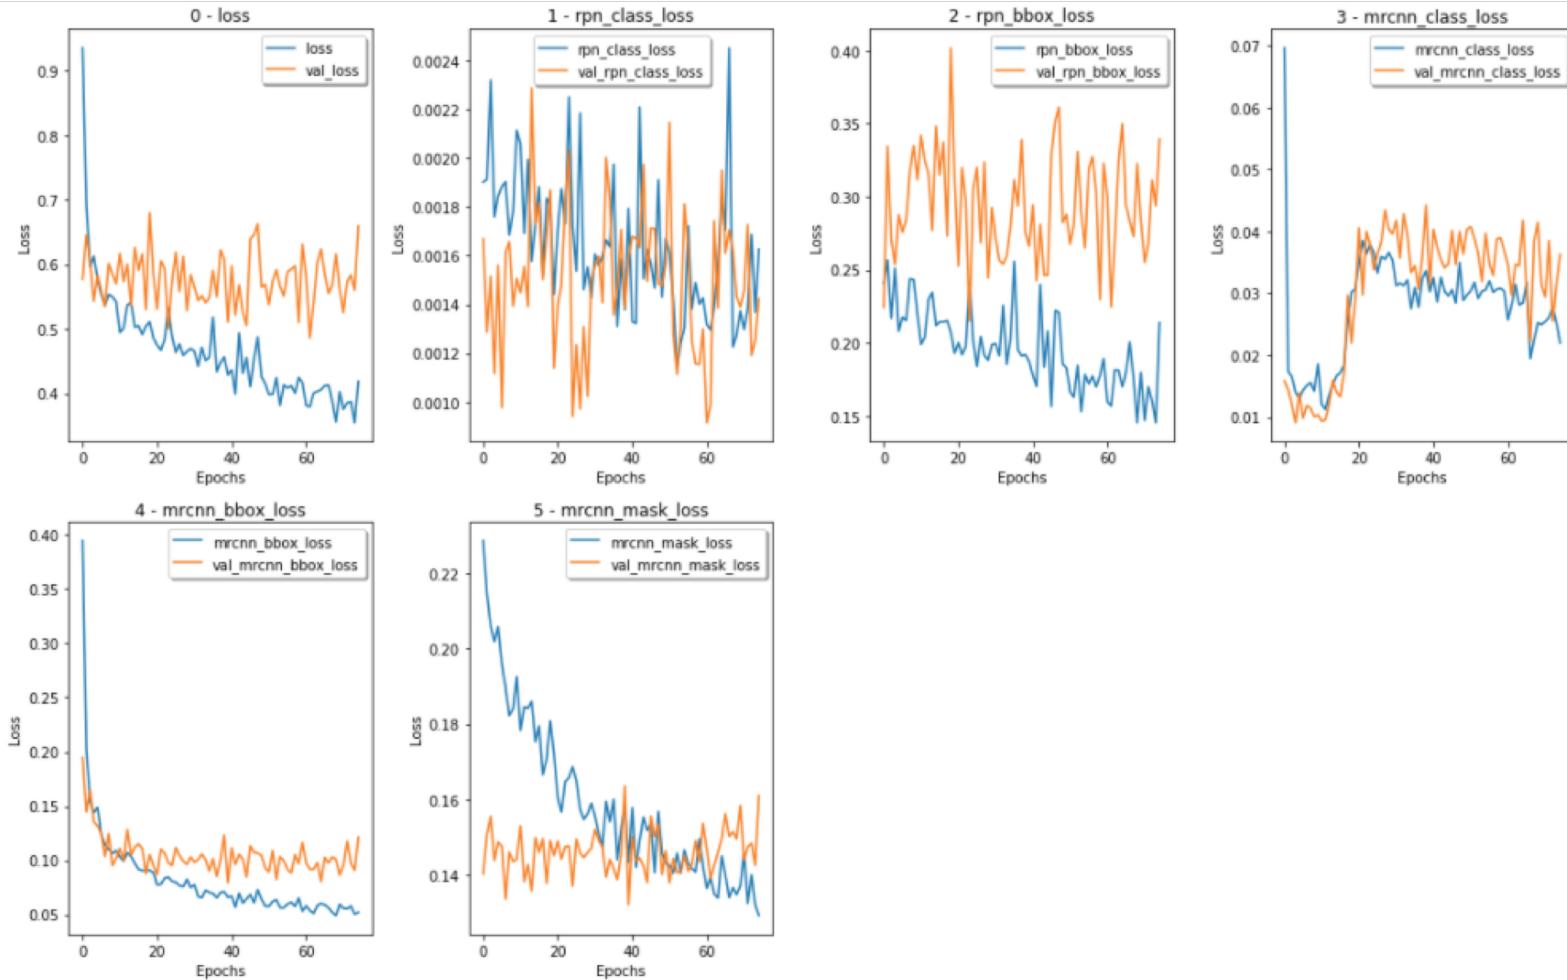
ACTIVATION	relu	ROI_POSITIVE_RATIO	0.33
ALL_CLASS NAMES	['BG', 'wolf']	RPN_ANCHOR RATIOS	[0.5, 1, 2]
BACKBONE	resnet101	RPN_ANCHOR SCALES	(32, 64, 128, 256, 512)
BACKBONE_STRIDES	[4, 8, 16, 32, 64]	RPN_ANCHOR STRIDE	1
BATCH_SIZE	2	RPN_BBOX_STD_DEV	[0.1 0.1 0.2 0.2]
BBOX_STD_DEV	[0.1 0.1 0.2 0.2]	RPN_NMS_THRESHOLD	0.7
CLASS NAMES	['wolf']	RPN_TRAIN_ANCHORS_PER_IMAGE	256
DETECTION_MAX_INSTANCES	100	STEPS_PER_EPOCH	100
DETECTION_MIN_CONFIDENCE	0.9	TRAINING_VERBOSE	1
DETECTION_NMS_THRESHOLD	0.3	TRAIN_BN	False
GPU_COUNT	1	TRAIN_ROIS_PER_IMAGE	200
GRADIENT_CLIP_NORM	5.0	USE_MINI_MASK	True
IMAGES_PER_GPU	2	USE_RPN_ROIS	True
LEARNING_RATE	0.001	VALIDATION_STEPS	50
LOSS_WEIGHTS	{'rpn_class_loss': 1.0, 'rpn_bbox_loss': 1.0, 'mrcnn_class_loss': 1.0, 'mrcnn_bbox_loss': 1.0, 'mrcnn_mask_loss': 1.0}	WEIGHT_DECAY	0.0001
MASK_POOL_SIZE	14		
MASK_SHAPE	[28, 28]		
MAX_GT_INSTANCES	100		
MEAN_PIXEL	[123.7 116.8 103.9]		
MINI_MASK_SHAPE	(56, 56)		
NAME	wolf		
NUM_CLASSES	2		
POOL_SIZE	7		
POST_NMS_ROIS_INFERENCE	1000		
POST_NMS_ROIS_TRAINING	2000		

# Example Training Output

```
Epoch 61/75100/100 [=====] - 67s 674ms/step - loss: 0.1700 - rpn_class_loss: 0.0018 -  
rpn_bbox_loss: 0.0581 - mrcnn_class_loss: 0.0085 - mrcnn_bbox_loss: 0.0302 - mrcnn_mask_loss: 0.0714 - val_loss:  
0.7379 - val_rpn_class_loss: 0.0055 - val_rpn_bbox_loss: 0.2347 - val_mrcnn_class_loss: 0.0662 -  
val_mrcnn_bbox_loss: 0.1594 - val_mrcnn_mask_loss: 0.2722  
Epoch 62/75100/100 [=====] - 67s 669ms/step - loss: 0.2485 - rpn_class_loss: 0.0029 -  
rpn_bbox_loss: 0.1213 - mrcnn_class_loss: 0.0101 - mrcnn_bbox_loss: 0.0364 - mrcnn_mask_loss: 0.0779 - val_loss:  
0.7327 - val_rpn_class_loss: 0.0047 - val_rpn_bbox_loss: 0.2240 - val_mrcnn_class_loss: 0.0743 -  
val_mrcnn_bbox_loss: 0.1654 - val_mrcnn_mask_loss: 0.2642  
Epoch 63/75100/100 [=====] - 67s 669ms/step - loss: 0.1893 - rpn_class_loss: 0.0023 -  
rpn_bbox_loss: 0.0708 - mrcnn_class_loss: 0.0091 - mrcnn_bbox_loss: 0.0341 - mrcnn_mask_loss: 0.0730 - val_loss:  
0.7406 - val_rpn_class_loss: 0.0039 - val_rpn_bbox_loss: 0.2385 - val_mrcnn_class_loss: 0.0793 -  
val_mrcnn_bbox_loss: 0.1763 - val_mrcnn_mask_loss: 0.2426  
Epoch 64/75100/100 [=====] - 67s 674ms/step - loss: 0.1678 - rpn_class_loss: 0.0014 -  
rpn_bbox_loss: 0.0549 - mrcnn_class_loss: 0.0072 - mrcnn_bbox_loss: 0.0314 - mrcnn_mask_loss: 0.0728 - val_loss:  
0.7117 - val_rpn_class_loss: 0.0049 - val_rpn_bbox_loss: 0.2295 - val_mrcnn_class_loss: 0.0665 -  
val_mrcnn_bbox_loss: 0.1654 - val_mrcnn_mask_loss: 0.2454  
Epoch 65/75100/100 [=====] - 68s 677ms/step - loss: 0.1649 - rpn_class_loss: 0.0016 -  
rpn_bbox_loss: 0.0548 - mrcnn_class_loss: 0.0050 - mrcnn_bbox_loss: 0.0317 - mrcnn_mask_loss: 0.0717 - val_loss:  
0.7583 - val_rpn_class_loss: 0.0041 - val_rpn_bbox_loss: 0.2411 - val_mrcnn_class_loss: 0.0660 -  
val_mrcnn_bbox_loss: 0.1838 - val_mrcnn_mask_loss: 0.2634
```

# Evaluating Loss Results

- I captured the history data return by Keras and plotted the loss functions of the different networks in the mode.



# Detection

- Example for running detection on a single image

```
# Set up paths and directory variable (not shown for readability)
```

```
from mrcnn import utils  
import mrcnn.model as modellib  
import csci_e89_project.det as det
```

```
# Path to trained weights
```

```
PROJECT_WEIGHTS_PATH = "models/wolf20180508T1124_mask_rcnn_wolf_0075.h5"
```

```
config = InferenceConfig()  
config.display()
```

```
model = modellib.MaskRCNN(mode="inference", model_dir=MODEL_DIR, config=config)  
model.load_weights(weights_path, by_name=True)
```

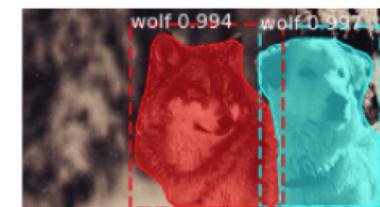
```
test_img_filename = TEST_IMAGE_DIR + 'wolf_pack.jpg'  
test_img = plt.imread(test_img_filename)
```

```
results = model.detect([test_img], verbose=1)
```

# Sample Result



# Other Results

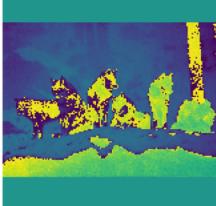
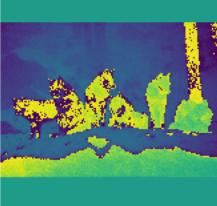
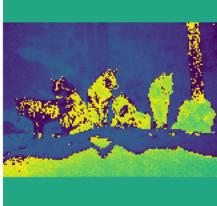
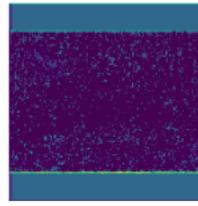
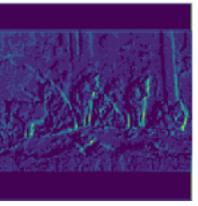
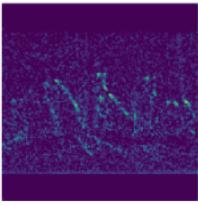
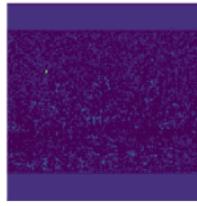
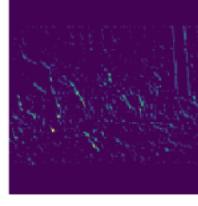
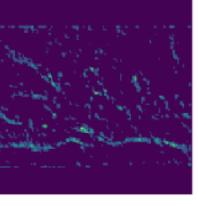
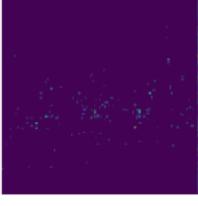
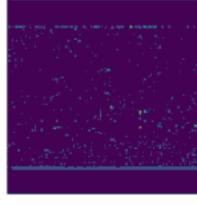
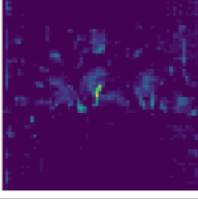
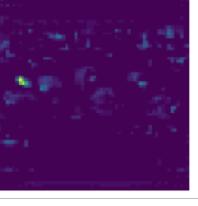
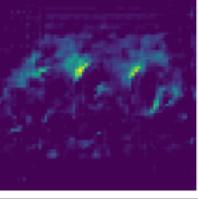
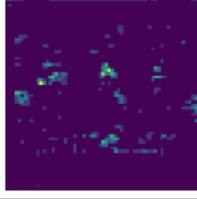


# Summary of Mask RCNN framework

## Process

- Stage 1: Regional Proposal Network (RPN)
  - RPN Targets
  - RPN Predictions
- Stage 2: For each proposal
  - Proposal Classification
  - Process detection
    - Filtering out background
    - Filter by score
    - Apply non-max suppression (NMS)
  - Apply Bounding Box Refinement
- Stage 3: Generating Masks

# Visualizing Activations

<u>Original</u>	<u>Layer</u>	<u>Activations</u>
	input_image	  
	res2c_out	   
	res3c_out	   
	res4w_out	   

# Summary of Results

- Mask RCNN framework is well written and easy to use. Some changes are required to
- Even with a small data set, a model trained with a single class showed impressive results.
- The overfitting could be attributed to training using 137 images, validation using 34. Additional images needed.
- The VIA annotation tools was simple and easy to use, and works well with small data sets.
- Detection time 0.250 seconds on GPU, 10 seconds on CPU

# Lessons Learned

- With small or limited data sets, leveraging pre-trained models (transfer learning) can still result in good performing models.
- The area of object detection and instance segmentation pulls together many concepts and networks to produce an end-end system.
- Interconnecting different modules (RPN, FPN, etc) is a powerful architectural feature.
- Not suitable for real-time. Detection on my Mac's CPU were approx 5-10 seconds.
- Other approaches such as single shot detectors and yolo are more promising for real time, but they do not support instance segmentation. Also, for accuracy, Faster R-CNN which is what Mask R-CNN is based is currently a good choice.
- VIA is not sufficient for complicated workflows or teams. Other tools should be considered.

# YouTube URLs, Last Page

- Two minute (short):
- 15 minutes (long):